


EC535 Introduction to Embedded Systems

- Teaming
- Blocking I/O
- Asynchronous Notification

Team: Pairs for
Assignments, HW, Labs

Last Time:

Metrics for Embedded Systems



Architecture-dependent Metrics

- Number of instructions/cycles, memory requirement, etc.
- More accurate, but also more expensive
 - May need to simulate or to use the hardware
- Useful to optimize a software for a target hardware
- Potential problems:
 - Simulation model may introduce some abstraction
 - e.g., a cycle-accurate simulator may not be able to model the processor bus and other peripherals
 - Simulated hardware may be different from the target

Instruction-accurate profiling

- SimIt-ARM instruction-accurate profiling:

```
$ ema -f nwfpe.bin small
```

```
The result is 4900
```

```
Total user time : 0.070 sec.
```

```
Total system time: 0.001 sec.
```

```
Simulation speed : 3.088e+07 inst/sec.
```

```
Total instructions : 2191774 (2M) including 7211 nullified
```

```
Total 4K memory pages allocated : 370
```

- $\text{Execution Time} = \text{Total Instructions} / \text{Simulation Speed}$

Cycle-accurate profiling

```
$ sima -f nwfpe.bin small
```

The result is 4900

Total icache reads: 2551776

Total icache read misses: 457

icache hit ratio: 99.982%

Total itlb reads: 2551819

Total itlb read misses: 25

itlb hit ratio: 99.999%

Total dcache writes: 369657

Total dcache write misses: 3654

Total dcache reads: 871393

Total dcache read misses: 207

dcache hit ratio: 99.689%

Total dtlb reads: 1241050

Total dtlb read misses: 26

dtlb hit ratio: 99.998%

Total biu accesses: 4315

biu activity: 3.438%

Total allocated OSMS : 2551819

Total retired OSMS : 2551817

Total cycles : 3131710

Equivalent time on 206.4MHz host: 0.0152 sec.

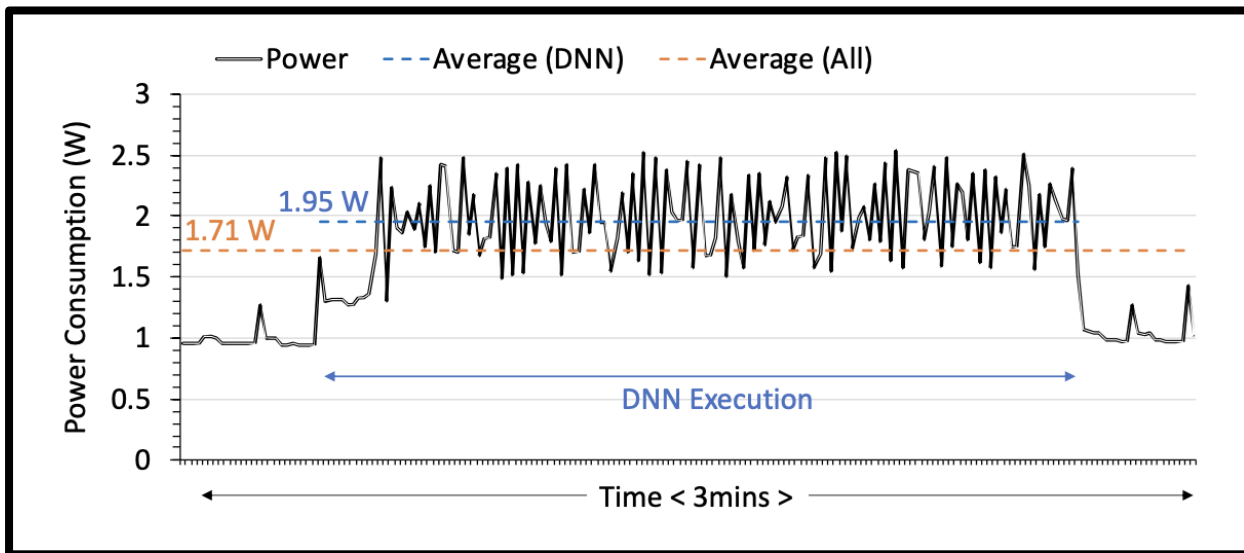
Static memory requirements

- Code and data space requirements using *size*:

```
$ arm-linux-size small
   text    data     bss      dec       hex filename
362479    4172     5140   371791   5ac4f small
```

- text: Code size
- data: Initialized data size
- bss: Non-initialized data size
- dec: Total bytes required in decimal
- hex: Total bytes required in hex

Other Metrics?



“We want to run as much as we can **on device**, because it delivers low latency...”

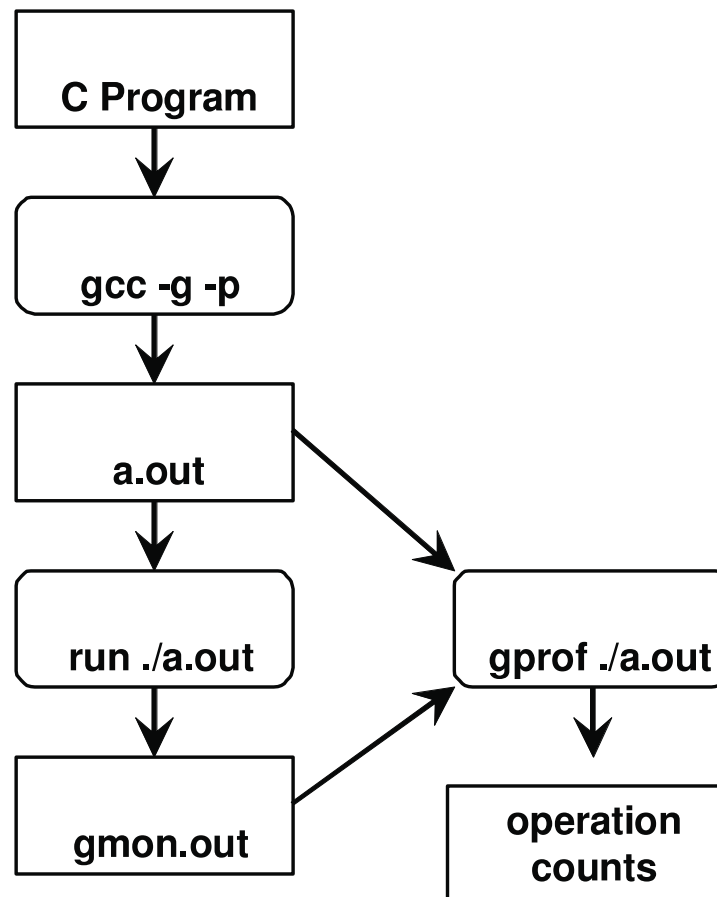


Profiling C Code

- Profiling:
 - Invasive: modify the program, i.e., code instrumentation
 - Non-invasive: statistic sampling of the program
- Profiles:
 - Flat profile
 - Call graph
 - Annotated sources
- Tools:
 - gprof
 - gcov
 - valgrind
 - oprofile

gprof for profiling C code

gcc inserts calls to a function mcount into prologue of each function



Profiling C code

■ small.c

```
#include "stdio.h"
int two(int limit) {
    int a, i;
    a = 0;
    for (i=0; i<limit; i++)
        a += i;
}

int one(int limit) {
    int i, a[50];

    for (i=0; i<limit; i++)
        a[i % 50] = i + two(i);
    return a[49];
}

int main() {
    int j, a;
    a = 0;
    for (j=0; j<1000; j++)
        a = a + one(j);
    printf("The result is %d\n", a);
    return 0;
}
```

gprof for profiling C code

- Enable profiling during compilation

```
> gcc -g -p small.c
```

- When run, the binary will create a file “gmon.out”, which can be analyzed by gprof

```
> ./a.out
```

```
> gprof a.out
```

- Some of the output looks like:

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
98.61	0.35	0.35	499500	0.71	0.71	two
1.39	0.36	0.01	1000	5.00	360.00	one

gprof for profiling C code

- two() is much more significant
- However, two() is called by one() unnecessarily 98% of the time

In-class exercise:

- Run gprof with small.c
- Optimize something outside the main function to speed up the program
- Run gprof again, observe the change in the flat profile
- Submit a zip file on GradeScope (inclass exercise 3) including:
 - New code in a file named small_new.c
 - Write names/usernames of people in the team as comments
 - The old and the new flat profile

```
#include "stdio.h"

int two(int limit) {
    int a, i;
    a = 0;
    for (i=0; i<limit; i++)
        a += i;
}

int one(int limit) {
    int i, a[50];

    for (i=0; i<limit; i++)
        a[i % 50] = i + two(i);
    return a[49];
}

int main() {
    int j, a;
    a = 0;
    for (j=0; j<1000; j++)
        a = a + one(j);
    printf("The result is %d\n", a);
    return 0;
}
```

gprof for profiling C code

```
int two(int limit)
{
    int a, i;
    a = 0;
    for (i=0; i<limit; i++)
        a += i;
}

int one(int limit)
{
    int i, a[50];

    for (i=49; i<limit; i+=50) {
        a[i % 50] = i + two(i);
    }
    return a[49];
}

int main()
{
    int j, a;
    a = 0;
    for (j=0; j<1000; j++)
        a = a + one(j);
    printf("The result is %d\n", a);
    return 0;
}
```

```
#include "stdio.h"
int two(int limit) {
    int a, i;
    a = 0;
    for (i=0; i<limit; i++)
        a += i;
}

int one(int limit) {
    int i, a[50];

    for (i=0; i<limit; i++)
        a[i % 50] = i + two(i);
    return a[49];
}

int main() {
    int j, a;
    a = 0;
    for (j=0; j<1000; j++)
        a = a + one(j);
    printf("The result is %d\n", a);
    return 0;
}
```


I/O-oriented programming in Linux

- I/O service is not always available
 - Bandwidth limitation
 - Data not yet ready when one wants to read
 - Old data not yet transmitted when one wants to write again
 - Asynchronous nature of I/O
 - Data arrival time is unpredictable



Simplistic Solution

- Nonblocking file operations on device file
 - If no data or device busy, read/write returns `-EAGAIN`
 - User process receives 0 as return value
 - If data available, read/write returns actual byte count transferred
 - User process receives the count
- Simplistic solution
 - When requesting data
 - User process keeps reading until read returns >0 .
 - When sending data
 - User process keeps writing until write returns <0 .



Better Solution



- Blocking file operations
 - If no data or device busy, read/write sleeps
 - Use the “wait queue” data structure
 - `wait_event_interruptible`
 - User process sleeps as a result
 - When I/O available, read/write wakes up
 - `wake_up_interruptible`
 - User process wakes up and resumes

Rules about sleeping



- Never sleep while running in atomic context.
 - What can go wrong?
 - Process holds a spinlock, other locks
 - Process disabled interrupts

Rules about sleeping



- Never sleep while running in atomic context.
 - What can go wrong?
 - Process holds a spinlock, other locks
 - Process disabled interrupts
- When process wakes up, it will not know what happened on the CPU while it was sleeping.
 - Make no assumptions based on earlier CPU states, check again.
- Make sure to verify who is going to wake the process up under what condition before implementing sleep.

Simple Sleep

- Wait queue
 - A linked list of sleeping processes
 - Wait events: interruptable, timeout, etc.
 - Use interruptable: can be interrupted by signals
 - Initialize the queue
 - Implement the calls to enter into wait queue and to finish waiting
 - Use `wake_up_interruptible`: restricts itself to processes that are on interruptable sleep (otherwise you may wake up all processes)
 - Implement a condition for `wake_up_interruptible`, e.g., a flag



Simple Sleep

- Wait queue
 - A linked list of sleeping processes
 - Wait events: interruptable, timeout, etc.
 - Use interruptable: can be interrupted by signals
 - Initialize the queue
 - Implement the calls to enter into wait queue and to finish waiting
 - Use `wake_up_interruptible`: restricts itself to processes that are on interruptable sleep (otherwise you may wake up all processes)
 - Implement a condition for `wake_up_interruptible`, e.g., a flag



Multiple Processes Sleeping?

Simple Sleep

- Wait queue
 - A linked list of sleeping processes
 - Wait events: interruptable, timeout, etc.
 - Use interruptable: can be interrupted by signals
 - Initialize the queue
 - Implement the calls to enter into wait queue and to finish waiting
 - Use `wake_up_interruptible`: restricts itself to processes that are on interruptable sleep (otherwise you may wake up all processes)
 - Implement a condition for `wake_up_interruptible`, e.g., a flag



Multiple Processes Sleeping?



Blocking vs. Non-blocking

- Potential problems with blocking operations
 - What if the program does not want to sleep?
 - e.g., need to handle other tasks.
 - Solution: can use multiple threads

Blocking vs. Non-blocking

- Potential problems with blocking operations
 - What if the program does not want to sleep?
 - e.g., need to handle other tasks.
 - Solution: can use multiple threads
- Separation of read/write buffers

Blocking vs. Non-blocking

- Potential problems with blocking operations
 - What if the program does not want to sleep?
 - e.g., need to handle other tasks.
 - Solution: can use multiple threads
- Separation of read/write buffers
- NONBLOCK flag: affects read, write, open operations
 - -EAGAIN: “try it again”
 - Useful for short operations that either succeed or fail quickly

Advanced sleep

- Set process state (running, interruptable/uninterruptable)
 - Interruptible/uninterruptible: two types of sleep
 - This step does not yield the processor yet
 - Discouraged → prone to bugs
- Check sleep condition and call the scheduler
 - `schedule()` will yield resources

Scull_pipe: a blocking I/O example
([\\$EC535/examples/ldd3_book](#))



Another Solution

- Asynchronous Notification
 - User process registers a signal handler function
 - Continues to do other tasks
 - When I/O available, kernel sends signal to user process
 - Signal handler is invoked

Setting I/O Modes

- Blocking/Async mode controlled by *f_flags* of *filp*
 - Changing bits in the flag changes device file behavior
- Fcntl system call can be used to read/write *f_flags*
 - `oldflag = fcntl(fd, F_GETFL);`
 - `fcntl(fd, F_SETFL, oldflag | FASYNC);`

Team - Shared Reading

- LDD3 Book
 - P. 147, Blocking I/O (Person 1)
 - P. 169, Asynchronous Notification (Person 2)

Team - Shared Reading

- LDD3 Book
 - P. 147, Blocking I/O (Person 1)
 - P. 169, Asynchronous Notification (Person 2)

Questions:

How do you wake up a sleeping process?

When a process receives a SIGIO, if two input files are available, how does the process know which input file has new input to offer?

What could go wrong if a process performing an atomic operation sleeps.

Advanced sleep

- Set process state (running, interruptable/uninterruptable)
 - Interruptible/uninterruptible: two types of sleep
 - This step does not yield the processor yet
 - Discouraged → prone to bugs
- Check sleep condition and call the scheduler
 - `schedule()` will yield resources

Scull_pipe: a blocking I/O example
([\\$EC535/examples/ldd3_book](#))